

Lecture 8: Metropolis Algorithms and Diagnostics

Merlise Clyde

September 21



Last Class: Normal Means Model

Data Model

$$Y_i \mid \mu_i, \sigma^2 \stackrel{iid}{\sim} (\mu_i, \sigma^2)$$

Means Model

$$\mu_i \mid \mu, \sigma_\mu^2 \stackrel{iid}{\sim} (\mu, \sigma_\mu^2)$$

Found marginal likelihood $\mathcal{L}(\mu, \sigma^2, \sigma_\mu^2)$ by integrating out μ_i with respect to g

$$\mathcal{L}(\mu, \sigma^2, \sigma_\mu^2) \propto \prod_{i=1}^n (\sigma^2 + \sigma_\mu^2)^{-1/2} \exp \left\{ -\frac{1}{2} \frac{(y_i - \mu)^2}{\sigma^2 + \sigma_\mu^2} \right\}$$

Posterior for $\theta = \mu, \sigma_\mu^2$ with $\sigma^2 = 1$

$$\pi(\theta \mid y) = \frac{\pi(\theta) \mathcal{L}(\theta)}{\int_{\Theta} \pi(\theta) \mathcal{L}(\theta) d\theta} = \frac{\pi(\theta) \mathcal{L}(\theta)}{m(y)}$$



Stochastic Sampling Intuition

- From a sampling perspective, we need to have a large sample or group of values, $\theta^{(1)}, \dots, \theta^{(S)}$ from $\pi(\theta \mid y)$ whose empirical distribution approximates $\pi(\theta \mid y)$.
- for any two sets A and B , we want

$$\frac{\frac{\#\theta^{(s)} \in A}{S}}{\frac{\#\theta^{(s)} \in B}{S}} = \frac{\#\theta^{(s)} \in A}{\#\theta^{(s)} \in B} \approx \frac{\pi(\theta \in A \mid y)}{\pi(\theta \in B \mid y)}$$

- Suppose we have a working group $\theta^{(1)}, \dots, \theta^{(s)}$ at iteration s , and need to add a new value $\theta^{(s+1)}$.
- Consider a candidate value θ^* that is *close* to $\theta^{(s)}$.
- Should we set $\theta^{(s+1)} = \theta^*$ or not?



Metropolis Ratio

look at the ratio

$$\begin{aligned} M &= \frac{\pi(\theta^* | y)}{\pi(\theta^{(s)} | y)} = \frac{\frac{p(y | \theta^*)\pi(\theta^*)}{p(y)}}{\frac{p(y | \theta^{(s)})\pi(\theta^{(s)})}{p(y)}} \\ &= \frac{p(y | \theta^*)\pi(\theta^*)}{p(y | \theta^{(s)})\pi(\theta^{(s)})} \end{aligned}$$

- does not depend on the marginal likelihood we don't know!



Metropolis algorithm

- If $M > 1$
 - Intuition: $\theta^{(s)}$ is already a part of the density we desire and the density at θ^* is even higher than the density at $\theta^{(s)}$.
 - Action: set $\theta^{(s+1)} = \theta^*$
- If $M < 1$,
 - Intuition: relative frequency of values in our group $\theta^{(1)}, \dots, \theta^{(s)}$ "equal" to θ^* should be $\approx M = \frac{\pi(\theta^* | y)}{\pi(\theta^{(s)} | y)}$.
 - For every $\theta^{(s)}$, include only a fraction of an instance of θ^* .
 - Action: set $\theta^{(s+1)} = \theta^*$ with probability M and $\theta^{(s+1)} = \theta^{(s)}$ with probability $1 - M$.



Proposal Distribution

- Where should the proposed value θ^* come from?
- Sample θ^* close to the current value $\theta^{(s)}$ using a **symmetric proposal distribution** $q(\theta^* \mid \theta^{(s)})$
- $q()$ is actually a "family of proposal distributions", indexed by the specific value of $\theta^{(s)}$.
- Here, symmetric means that $q(\theta^* \mid \theta^{(s)}) = q(\theta^{(s)} \mid \theta^*)$.
- Common choice

$$N(\theta^*; \theta^{(s)}, \delta^2 \Sigma)$$

with Σ based on the approximate $\text{Cov}(\theta \mid y)$ and
 $\delta = 2.44/\text{dim}(\theta)$ or tune

$$\text{Unif}(\theta^*; \theta^{(s)} - \delta, \theta^{(s)} + \delta)$$



Metropolis Algorithm Recap

- The algorithm proceeds as follows:

1. Given $\theta^{(1)}, \dots, \theta^{(s)}$, generate a *candidate* value $\theta^* \sim q(\theta^* \mid \theta^{(s)})$.

2. Compute the acceptance ratio

$$M = \frac{\pi(\theta^* \mid y)}{\pi(\theta^{(s)} \mid y)} = \frac{p(y \mid \theta^*)\pi(\theta^*)}{p(y \mid \theta^{(s)})\pi(\theta^{(s)})}.$$

3. Set

$$\theta^{(s+1)} = \begin{cases} \theta^* & \text{with probability } \min(M, 1) \\ \theta^{(s)} & \text{with probability } 1 - \min(M, 1) \end{cases}$$

equivalent to sampling $u \sim U(0, 1)$ independently and setting

$$\theta^{(s+1)} = \begin{cases} \theta^* & \text{if } u < M \\ \theta^{(s)} & \text{if otherwise} \end{cases}.$$



Metropolis algorithm

- Once we obtain the samples, then we are back to using Monte Carlo approximations for quantities of interest!
- we can approximate posterior means, quantiles, and other quantities of interest using the empirical distribution of our sampled values.
- easy to compute the posterior distribution of nonlinear functions of parameters!

$$\psi^{(s)} = g(\theta^{(s)})$$

- some posterior summaries are hard to calculate based on samples $\{\theta^{(s)}\}$
 - mode/MAP (at least for continuous)
 - marginal likelihood $m(y) = \int \pi(\theta)p(y | \theta) d\theta$



Notes

- The Metropolis chain ALWAYS moves to the proposed θ^* at iteration $s + 1$ if θ^* has higher target density than the current $\theta^{(s)}$.
- Sometimes, it also moves to a θ^* value with lower density in proportion to the density value itself.
- This leads to a random, Markov process that naturally explores the space according to the probability defined by $\pi(\theta \mid y)$, and hence generates a sequence that, while dependent, eventually represents draws from $\pi(\theta \mid y)$ (stationary distribution of the Markov Chain).



Convergence

We will not cover the convergence theory behind Metropolis chains in detail, but ...

- The Markov process generated under this procedure is **ergodic** (irreducible and aperiodic) and has a unique limiting distribution (stationary distribution)
 - ergodicity means that the chain can move anywhere at each step, which is ensured, if $q(\theta^* \mid \theta^{(s)}) > 0$ everywhere!
- By construction, Metropolis chains are **reversible**, so that $\pi(\theta \mid y)$ is the stationary distribution
 - Think of reversibility as being equivalent to symmetry of the joint density of two consecutive $\theta^{(s)}$ and $\theta^{(s+1)}$ in the stationary process (which we get by using a symmetric proposal distribution)
 - detailed balance



Example


Priors with $\sigma^2 = 1$:

$$p(\mu) \propto 1$$

- Use a **Cauchy**(0, 1) prior on σ_μ independent of μ and
- Symmetric proposal for μ and σ_τ ?
- Try independent normals $\frac{2.44^2}{d} \text{Cov}(\theta)$ where d is the dimension of θ ($d = 2$)



First 200 Samples



- Overall Acceptance probability is 0.6 out of 10^4 samples
- Goal is around 0.44 in 1 dimension to 0.23 in higher dimensions



Tuning

- Sampled values are correlated
- Correlation between samples can be adjusted by selecting an optimal δ (i.e., spread of the distribution) in the proposal distribution
- δ too small leads to $M \approx 1$ for most proposed values, a high acceptance rate, but very small moves, leading to highly correlated chain.
- δ too large can get "stuck" because θ^* may be very far away from high density regions, leading to a very low acceptance rate and again high correlation in the Markov chain.
- Burn-in and thinning can help!



Burn-in

- Convergence occurs regardless of our starting point (in theory), so we can usually pick any reasonable values in the parameter space as a starting point.
- May take a long time to reach high density regions
- Over representation of low density samples given finite iterations
- Generally, we throw out a certain number of the first draws, known as the **burn-in**, as an attempt to make our draws closer to the stationary distribution and less dependent on any single set of starting values.
- However, we don't know exactly when convergence occurs, so it is not always clear how much burn-in we would need.
- If you run long enough you should not need to discard any samples! (ergodicity)



Convergence diagnostics

- Diagnostics available to help decide on number of burn-in & collected samples.
- **Note:** no definitive tests of convergence but you should do as many diagnostics as you can, on all parameters in your model.
- With "experience", visual inspection of trace plots perhaps most useful approach.
- There are a number of useful automated tests in R.
- **CAUTION:** diagnostics cannot guarantee that a chain has converged, but they can indicate it has not converged.



Diagnostics in R

- The most popular package for MCMC diagnostics in R is coda.
- coda uses a special MCMC format so you must always convert your posterior matrix into an MCMC object.
- For the example, we have the following in R.

```
#library(coda)  
theta.mcmc <- mcmc(theta,start=1) #no burn-in (simple problem!
```



Diagnostics in R

```
summary(theta.mcmc)
```

```
##
## Iterations = 1:10000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean      SD Naive SE Time-series SE
## mu          -0.07977 0.1046 0.001046      0.002839
## sigma_mu     0.17550 0.1273 0.001273      0.004397
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%   97.5%
## mu          -0.283420 -0.1508 -0.08193 -0.00848 0.1337
## sigma_mu     0.007995  0.0758  0.15024  0.25228 0.4693
```



The naive SE is the **standard error of the mean**, which captures simulation error of the mean rather than the posterior uncertainty.

Effective sample size

- The **effective sample size** translates the number of MCMC samples S into an equivalent number of independent samples.
- It is defined as

$$\text{ESS} = \frac{S}{1 + 2 \sum_k \rho_k},$$

where S is the sample size and ρ_k is the lag k autocorrelation.

- For our data, we have

```
effectiveSize(theta.mcmc)
```

```
##           mu  sigma_mu  
## 1356.6495  838.2613
```


- So our 10,000 samples are equivalent to 1356.6 independent samples for μ and 838.3 independent samples for σ_μ .




Trace plot for mean

```
plot(theta.mcmc[, "mu"])
```

Trace of var1



Density of var1




N = 10000 Bandwidth = 0.01757




Trace plot for σ_μ

```
plot(theta.mcmc[, "sigma_mu"])
```

Trace of var1



Density of var1



N = 10000 Bandwidth = 0.02139



OK (be careful of scaling in plots!)

Autocorrelation

- Another way to evaluate convergence is to look at the autocorrelation between draws of our Markov chain.
- The lag k autocorrelation, ρ_k , is the correlation between each draw and its k th lag, defined as


$$\rho_k = \frac{\sum_{s=1}^{S-k} (\theta_s - \bar{\theta})(\theta_{s+k} - \bar{\theta})}{\sum_{s=1}^{S-k} (\theta_s - \bar{\theta})^2}.$$

- We expect the autocorrelation to decrease as k increases.
- If autocorrelation remains high as k increases, we have slow mixing due to the inability of the sampler to move around the space well.



Autocorrelation for mean


```
autocorr.plot(theta.mcmc[, "mu"])
```



So-So

Autocorrelation for variance


```
autocorr.plot(theta.mcmc[, "sigma_mu"])
```



worse

Gelman-Rubin

Gelman & Rubin suggested a diagnostic R based on taking separate chains with dispersed initial values to test convergence



Gelman-Rubin Diagnostic

- Run $m > 2$ chains of length $2S$ from overdispersed starting values.
- Discard the first S draws in each chain.
- Calculate the pooled within-chain variance W and between-chain variance B .

$$R = \frac{\frac{S-1}{S}W + \frac{1}{S}B}{W}$$

- numerator and denominator are both unbiased estimates of the variance if the two chains have converged
 - otherwise W is an underestimate (hasn't explored enough)
 - numerator will overestimate as B is too large (overdispersed starting points)
- As $S \rightarrow \infty$ and $B \rightarrow 0$, $R \rightarrow 1$
- version in R is slightly different



Gelman-Rubin Diagnostic

```
theta.mcmc = mcmc.list(mcmc(theta1, start=5000), mcmc(theta2, start=5000))
gelman.diag(theta.mcmc)
```

```
## Potential scale reduction factors:
```

```
##
```

```
##           Point est. Upper C.I.
```

```
## mu                1          1
```

```
## sigma_mu          1          1
```

```
##
```

```
## Multivariate psrf
```

```
##
```

```
## 1
```

- Values of $R > 1.1$ suggest lack of convergence
- Looks OK

See also `gelman.plot`



Geweke statistic

- Geweke proposed taking two non-overlapping parts of a single Markov chain (usually the first 10% and the last 50%) and comparing the mean of both parts, using a difference of means test
- The null hypothesis would be that the two parts of the chain are from the same distribution.
- The test statistic is a z-score with standard errors adjusted for autocorrelation, and if the p-value is significant for a variable, you need more draws.



Geweke Diagnostic

- The output is the z-score itself (not the p-value).

```
geweke.diag(theta.mcmc)
```

```
## [[1]]  
##  
## Fraction in 1st window = 0.1  
## Fraction in 2nd window = 0.5  
##  
##      mu sigma_mu  
## -0.7779  0.7491  
##  
##  
## [[2]]  
##  
## Fraction in 1st window = 0.1  
## Fraction in 2nd window = 0.5  
##  
##      mu sigma_mu  
##  0.4454  0.6377
```



Practical advice on diagnostics

- There are more tests we can use: Raftery and Lewis diagnostic, Heidelberger and Welch, etc.
- The Gelman-Rubin approach is quite appealing in using multiple chains
- Geweke (and Heidelberger and Welch) sometimes reject even when the trace plots look good.
- Overly sensitive to minor departures from stationarity that do not impact inferences.
- Most common method of assessing convergence is visual examination of trace plots.



Improving

- more iterations and multiple chains
- thinning to reduce correlations and increase ESS
- change the proposal distribution q

